Lab Guide 6

Optimizing performance & task parallelism

Objective:
- performance measurement to optimize parallel execution
- introduce the task concept

Introduction

This lab session introduces the `perf` tool to profile the execution of an application code on the PU.
The exercises aim to introduce performance optimizations on shared memory programming.
Copy the file `/share/cpar/PL06_Codigo` to your `home`, which is a program that sorts a vector twice.
Run the application with `perf record ./a.out` to sample the execution data at fixed time intervals.
Use `perf report` to generate a [flat] view profile with the application hotspots.
For better accuracy, the code (C program) should be compiled with `-g -fno-omit-frame-pointer`.

Exercise 1 - `Critical, atomic and Reduction` directives

Consider the following OpenMP program used in the previous lab session:

```c
#include<omp.h>
#include<stdio.h>

double f( double a ) {
    return (4.0 / (1.0 + a*a));
}
double pi = 3.141592653589793;
int main() {
    double mypi = 0;
    int n = 10000000; // number of points to compute
    float h = 1.0 / n;
    #pragma omp parallel for reduction(+:mypi)
    for(int i=0; i<n; i++) {
        mypi = mypi + f(i*h);
    }
    mypi = mypi * h;
    printf(" pi = %.10f \n", mypi);
}
```

a) Measure the code scalability by comparing `critical`, `atomic` and `reduction` directives to avoid the data race in the shared variable (`mypi`).
To obtain the execution times of your code use `batch --partition=cpar time.sh`

b) Analyse the time overhead of the `critical` directive using the `perf` tool.
Adapt the script `perf1.sh` for your case.

Exercise 2 – Develop an OpenMP code to implement the parallel execution of the QuickSort

```
void quickSort(float* arr, int size)
{
        int i = 0, j = size;
        /* PARTITION PART */
        partition(arr, &i, &j);

        if (0 < j){ quickSort_internel(arr, 0, j);}
        if (i< size){ quickSort_internel(arr, i, size);}
}
```

a) Run the original code (`batch --partition=cpar run.sh`) and explain why one of runs takes longer to execute.

b) Analyse the algorithm and suggest one parallel implementation based on `omp task`.

c) Analyse the algorithm complexity of the sequential and parallel fractions, knowing that the average algorithm complexity of this sorting is $N\log_2 N$.
Estimate the maximum parallelization gain for the problem size of 2048, with 2 tasks.

d) Run the code with 2, 4 and 8 threads, measure the scalability and explain the results.
Run the program with `sbath --partition=cpar run2.sh`.
Use `perf` to obtain the execution profile (remove `-fopenmp` from the compilation and run the program with `sbath --partition=cpar perf2.sh`).

e) Tuning: modify the parallelization approach to create tasks by the recursive call.
Execute the program with 1, 2 and 4 threads and explain the results.

f) Tuning: remove `task` creation when the size of the sub-problem is less or equal to 100 (parallelism cut-off).
Execute the program with 1, 2 and 4 threads and explain the results.
What is the best parallelism cut-off on this server?