




University of Minho
School of Engineering



Dados e Aprendizagem Automática

Artificial Neural Networks

DAA @ MEI-1º/MiEI-4º/MMC-1º – 1º Semestre
Bruno Fernandes, Filipe Gonçalves, Víctor Alves, Cesar Analide



What about Artificial Neural Networks?

3

Let's get back to our "real estate agent" problem and help our neighbor (the real estate agent) to predict housing prices for regions in the USA (dataset [here](#)).

We have already used Linear Regression... But now let's try using **Multilayer Perceptrons (MLPs)**, a class of **Artificial Neural Networks**!

To implement our first **Artificial Neural Network** we will use:



TensorFlow for Artificial Neural Networks

4

Why?

- Open-source software library for **high performance numerical computation**
- Strong support for **machine learning** and **deep learning**
- It has seen tremendous growth and popularity in the machine learning community

An open source machine learning library for research and production.

Companies using TensorFlow



Implementing a MLP - The Dataset

5

We already know this dataset (let's drop the *address* - we don't need it)! We will work a **regression problem**!

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Check the data (again...)

```
USAhousing = pd.read_csv('USA_Housing.csv')
USAhousing.drop('Address', axis=1, inplace=True)
USAhousing.head()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05


Implementing a MLP - Some Data Viz.

6

We have already explored it before so... Let's move on!


```
sns.distplot(USAhousing['Price'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x2972eb2ea90>
```



```
sns.heatmap(USAhousing.corr())
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x2972f437240>
```



Implementing a MLP - More imports!

7

To implement **our first MLP** we will need some more libraries! Let's import them all at once! You'll need to install **TensorFlow** - use the Navigator or the Prompt:

```
conda install -c conda-forge tensorflow
```

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, BatchNormalization
from tensorflow.keras.wrappers.scikit_learn import KerasRegressor
from sklearn.model_selection import GridSearchCV, KFold, train_test_split
from sklearn.preprocessing import MinMaxScaler
```

```
RANDOM_SEED = 2021
```

```
print("TensorFlow version:", tf.__version__)
```

```
TensorFlow version: 2.3.0
```

Implementing a MLP - Scaling the data

8

We can then define our **X** and **y**, as usual!

```
X = USAhousing.drop('Price', axis=1)
y = USAhousing[['Price']]
```

Artificial neural networks are picky - **they prefer scaled data**! Let's scale the data to be in the interval [0, 1].

```
# Let's scale the features between [0-1]
scaler_X = MinMaxScaler(feature_range=(0, 1)).fit(X)
scaler_y = MinMaxScaler(feature_range=(0, 1)).fit(y)
X_scaled = pd.DataFrame(scaler_X.transform(X[X.columns]), columns=X.columns)
y_scaled = pd.DataFrame(scaler_y.transform(y[y.columns]), columns=y.columns)
```


Implementing a MLP - Scaling the data

9

Let's just make **some checks!**

Everything is OK!

```
x.head()
```

	Avg. Area Income	Avg. Area House Age
0	79545.458574	5.682861
1	79248.642455	6.002900
2	61287.067179	5.865890
3	63345.240046	7.188236
4	59982.197226	5.040555

```
x_scaled.head()
```

	Avg. Area Income	Avg. Area House Age
0	0.686822	0.441986
1	0.683521	0.488538
2	0.483737	0.468609
3	0.506630	0.660956
4	0.469223	0.348556

```
y.head()
```

	Price
0	1.059034e+06
1	1.505891e+06
2	1.058988e+06
3	1.260617e+06
4	6.309435e+05

```
y_scaled.head()
```

	Price
0	0.425210
1	0.607369
2	0.425192
3	0.507384
4	0.250702

Implementing a MLP - Let's build the model!

10

Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

```
def build_model(activation='relu', learning_rate=0.01):  
    #Create a sequential model (with three layers - last one is the output)  
    model = Sequential()  
    model.add(Dense(16, input_dim=5, activation=activation))  
    model.add(Dense(8, activation=activation))  
    model.add(Dense(1, activation='relu'))  
  
    #Compile the model  
    #Define the loss function, the optimizer and metrics to be used  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```

Implementing a MLP - Let's build the model!

11

Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

```
def build_model(activation='relu', learning_rate=0.01):  
    #Create a sequential model (with three layers - last one is the output)  
    model = Sequential()  
    model.add(Dense(16, input_dim=5, activation=activation))  
    model.add(Dense(8, activation=activation))  
    model.add(Dense(1, activation='relu'))  
  
    #Compile the model  
    #Define the loss function, the optimizer and metrics to be used  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```

Layers stacked one
over the other!

Implementing a MLP - Let's build the model!

12

Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

```
def build_model(activation='relu', learning_rate=0.01):  
    #Create a sequential model (with three layers - last one is the output)  
    model = Sequential()  
    model.add(Dense(16, input_dim=5, activation=activation))  
    model.add(Dense(8, activation=activation))  
    model.add(Dense(1, activation='relu'))  
  
    #Compile the model  
    #Define the loss function, the optimizer and metrics to be used  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```

Layers stacked one
over the other!

Number of neurons
in each layer (last
one is the output).

Implementing a MLP - Let's build the model!

13

Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

```
def build_model(activation='relu', learning_rate=0.01):  
    #Create a sequential model (with three layers - last one is the output)  
    model = Sequential()  
    model.add(Dense(16, input_dim=5, activation=activation))  
    model.add(Dense(8, activation=activation))  
    model.add(Dense(1, activation='relu'))  
  
    #Compile the model  
    #Define the loss function, the optimizer and metrics to be used  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```

Layers stacked one
over the other!

Number of neurons
in each layer (last
one is the output).

Number of input
features.

Implementing a MLP - Let's build the model!

14

Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

```
def build_model(activation='relu', learning_rate=0.01):  
    #Create a sequential model (with three layers - last one is the output)  
    model = Sequential()  
    model.add(Dense(16, input_dim=5, activation=activation))  
    model.add(Dense(8, activation=activation))  
    model.add(Dense(1, activation='relu'))  
  
    #Compile the model  
    #Define the loss function, the optimizer and metrics to be used  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```

Layers stacked one over the other!

Number of neurons in each layer (last one is the output).

Number of input features.


Activation function (here, as an argument to the *build_model* function).

Implementing a MLP - Let's build the model!

15

Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

```
def build_model(activation='relu', learning_rate=0.01):  
    #Create a sequential model (with three layers - last one is the output)  
    model = Sequential()  
    model.add(Dense(16, input_dim=5, activation=activation))  
    model.add(Dense(8, activation=activation))  
    model.add(Dense(1, activation='relu'))  
  
    #Compile the model  
    #Define the loss function, the optimizer and metrics to be used  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```



After building the stacked MLP, we need to **compile the model** by setting the **loss function** (MSE as we are solving a regression problem), the **optimizer** (which implements the gradient descent and updates the weights), and a set of **metrics** (to further understand the performance of the model - not used when backpropagating the error).

Implementing a MLP - Train/Test data

16

Let's use the `train_test_split` API to **hold-out some data for testing!** We will use **cross validation over the training data!**

Test and training data

[illegible]

Implementing a MLP - Tuning the network

17

We want to **find the best possible MLP** to solve our problem so... Let's **tune it** (at least, some hyperparameters)!

We must first define a dictionary of {key -> *list of values*}. For now, we will only **tune** the **activation function** and the **learning rate** used by the optimizer (trying two values for each hyperparameter - so, **we will fit 4 candidate models**). Note that these are arguments of the *build_model()* function.

```
TUNING_DICT = {  
    'activation' :    ['relu', 'sigmoid'],  
    'learning_rate' : [0.01, 0.001]  
}
```

Implementing a MLP - Tuning the network

18

Now, let's use the **KFold API** with $k=5$, the **KerasRegressor API** to point to our *build_model()* function, and use the **GridSearchCV API** as we've done before!

```
kf = KFold(n_splits=5, shuffle=True, random_state=RANDOM_SEED)

model = KerasRegressor(build_fn=build_model, epochs=20, batch_size=32)
grid_search = GridSearchCV(estimator = model,
                           param_grid = TUNING_DICT,
                           cv = kf,
                           scoring = 'neg_mean_absolute_error',
                           refit = 'True',
                           verbose = 1)
grid_search.fit(X_train, y_train, validation_split=0.2, verbose=1)
```


Implementing a MLP - Tuning the network

19

Now, let's use the **KFold API** with $k=5$, the **KerasRegressor API** to point to our *build_model()* function, and use the **GridSearchCV API** as we've done before!

```
kf = KFold(n_splits=5, shuffle=True, random_state=RANDOM_SEED)

model = KerasRegressor(build_fn=build_model, epochs=20, batch_size=32)
grid_search = GridSearchCV(estimator = model,
                           param_grid = TUNING_DICT,
                           cv = kf,
                           scoring = 'neg_mean_absolute_error',
                           refit = 'True',
                           verbose = 1)
grid_search.fit(X_train, y_train, validation_split=0.2, verbose=1)
```



We must pass our **build_model** function, define the number of **epochs** (*number of passes of the entire training dataset*) and the **batch size** (*number of training samples in one forward/backward pass*). The KerasRegressor API will **return an instance of our MLP**.

Implementing a MLP - Tuning the network

20

Now, let's use the **KFold API** with $k=5$, the **KerasRegressor API** to point to our *build_model()* function, and use the **GridSearchCV API** as we've done before!

```
kf = KFold(n_splits=5, shuffle=True, random_state=RANDOM_SEED)

model = KerasRegressor(build_fn=build_model, epochs=20, batch_size=32)
grid_search = GridSearchCV(estimator = model,
                           param_grid = TUNING_DICT,
                           cv = kf,
                           scoring = 'neg_mean_absolute_error',
                           refit = 'True',
                           verbose = 1)
grid_search.fit(X_train, y_train, validation_split=0.2, verbose=1)
```

Fraction of the training data to be used as **validation data**. The model will set apart this fraction of the training data, **will not train on it**, and **will evaluate the loss and the model metrics** on this data at the end of each epoch.

We must pass our **build_model** function, define the number of **epochs** (*number of passes of the entire training dataset*) and the **batch size** (*number of training samples in one forward/backward pass*). The KerasRegressor API will **return an instance of our MLP**.

Implementing a MLP - Tuning the network

21

Now, let's use the **KFold API** with $k=5$, the **KerasRegressor API** to point to our *build_model()* function, and use the **GridSearchCV API** as we've done before!

```
kf = KFold(n_splits=5, shuffle=True, random_state=RANDOM_SEED)

model = KerasRegressor(build_fn=build_model, epochs=20, batch_size=32)
grid_search = GridSearchCV(estimator = model,
                           param_grid = TUNING_DICT,
                           cv = kf,
                           scoring = 'neg_mean_absolute_error',
                           refit = 'True',
                           verbose = 1)
grid_search.fit(X_train, y_train, validation_split=0.2, verbose=1)

100/100 [=====] - 0s 522us/step - loss: 0.0342 - mae: 0.0342 - mse: 0.0018 - val_loss: 0.0343 - val_
mae: 0.0341 - val_mse: 0.0018
Epoch 20/20
100/100 [=====] - 0s 520us/step - loss: 0.0335 - mae: 0.0335 - mse: 0.0017 - val_loss: 0.0343 - val_
mae: 0.0343 - val_mse: 0.0018

GridSearchCV(cv=KFold(n_splits=5, random_state=2021, shuffle=True),
             estimator=<tensorflow.python.keras.wrappers.scikit_learn.KerasRegressor object at 0x000001B5463DF748>,
             param_grid={'activation': ['relu', 'sigmoid'],
                          'learning_rate': [0.01, 0.001]},
             refit='True', scoring='neg_mean_absolute_error', verbose=1)
```

Implementing a MLP - Cross Validation results

22

We can now **analyze the performance** of our MLP!

```
#summarize results
print("Best: %f using %s" % (grid_search.best_score_, grid_search.best_params_))
means = grid_search.cv_results_['mean_test_score']
stds = grid_search.cv_results_['std_test_score']
params = grid_search.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

Best: -0.033692 using {'activation': 'relu', 'learning_rate': 0.001}

Our best model!

-0.130550 (0.184629) with: {'activation': 'relu', 'learning_rate': 0.01}
-0.033692 (0.000963) with: {'activation': 'relu', 'learning_rate': 0.001}
-0.403154 (0.182433) with: {'activation': 'sigmoid', 'learning_rate': 0.01}
-0.221557 (0.222091) with: {'activation': 'sigmoid', 'learning_rate': 0.001}

Implementing a MLP - Overfitting Analysis

23

We can also try to **understand if our model overfitted!**


```
#Our best model (remember we set refit=True?)  
best_mlp_model = grid_search.best_estimator_
```

```
#Did the model overfit?  
def plot_learning_curve(history, metric='neg_mean_absolute_error'):  
    plt.figure(figsize=(8,4))  
    plt.title('Training Loss vs Validation Loss')  
    plt.plot(history.epoch, history.history['loss'], label='train')  
    plt.plot(history.epoch, history.history['val_loss'], label='val')  
    plt.ylabel('Training ' + metric)  
    plt.xlabel('Epochs')  
    plt.legend()  
  
plot_learning_curve(best_mlp_model.model.history, metric='neg_mean_absolute_error')
```

Implementing a MLP - Overfitting Analysis

24

We can also try to **understand if our model overfitted!** I guess not!



Implementing a MLP - Performance Analysis

25

What about MLP predictions?

```
#Obtain predictions  
predictions = best_mlp_model.predict(X_test)  
predictions = predictions.reshape(predictions.shape[0], 1)  
predictions[:5]
```

```
array([[0.62522596],  
       [0.35357618],  
       [0.27740446],  
       [0.6624256 ],  
       [0.30916768]], dtype=float32)
```

Yap... **Predictions are scaled!!!** We can, however, use the *inverse_transform()* function to obtain the real values!

Implementing a MLP - Performance Analysis

26

What about MLP predictions?

```
#Obtain predictions
predictions = best_mlp_model.predict(X_test)
predictions = predictions.reshape(predictions.shape[0], 1)
predictions[:5]
```

```
array([[0.62522596],
       [0.35357618],
       [0.27740446],
       [0.6624256 ],
       [0.30916768]], dtype=float32)
```

Yap... **Predictions are scaled!!!** We can, however, use the *inverse_transform()* function to obtain the real values!

```
#And now let's unscale the model's predictions to see real prices!
predictions_unscaled = scaler_y.inverse_transform(predictions)
predictions_unscaled[:5]
```

```
array([[1511417.2 ],
       [ 849883.9 ],
       [ 670154.06],
       [1599453.5 ],
       [ 734485.9 ]], dtype=float32)
```

Implementing a MLP - Performance Analysis

27

Let's do the same for the `y_test` data (you can check that you will obtain the original values)!

```
#Let's unscale y_test to get the original values  
y_test_unscaled = scaler_y.inverse_transform(y_test)  
y_test_unscaled[:5]
```

```
array([[1409892.08977612],  
       [ 889385.90158426],  
       [ 635429.23051901],  
       [1613414.23305073],  
       [ 774491.65328819]])
```

```
#And now let's unscale the model's predictions to see real prices!  
predictions_unscaled = scaler_y.inverse_transform(predictions)  
predictions_unscaled[:5]
```

```
array([[1511417.2 ],  
       [ 849883.9 ],  
       [ 670154.06],  
       [1599453.5 ],  
       [ 734485.9 ]], dtype=float32)
```

Not that bad!!

Implementing a MLP - Performance Analysis


28

Some more metrics (scaled version)...

```
from sklearn import metrics
print('MAE:', metrics.mean_absolute_error(y_test, predictions))
print('MSE:', metrics.mean_squared_error(y_test, predictions))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

```
MAE: 0.033830947270123486
MSE: 0.0017417010324501016
RMSE: 0.041733691814289584
```

```
plt.scatter(y_test, predictions)
```



Implementing a MLP - Performance Analysis

29

Comparing **real** and **predicted values**...


```
#Visualising the actual and predicted result
def real_predicted_viz(limit):
    plt.figure(figsize=(14,6))
    plt.plot(y_test_unscaled[:limit], color = 'green', label = 'Actual')
    plt.plot(predictions_unscaled[:limit], color = 'red', label = 'Predicted')
    plt.grid(alpha = 0.3)
    plt.xlabel('Houses')
    plt.ylabel('Price')
    plt.title('Real vs Predicted')
    plt.legend()
    plt.show()

#Let's limit to 200 comparasions for better visualization
real_predicted_viz(200)
```

Implementing a MLP - Performance Analysis

30

Comparing **real** and **predicted values**...



Hands On

31

Hands On

The screenshot displays the Spyder Python IDE interface. The main editor window shows a Python script for a Mean-Shift algorithm. The script defines a `Mean_Shift` class with an `__init__` method to set parameters like `radius` and `radius_normalize_step`, and a `fit` method that iteratively finds centroids and updates weights based on the distance of data points. The script is located at `C:\data\PythonWorkspace\dev\meanshift_algorithm.py`.

The Variable explorer on the right shows the state of the program's variables:

Name	Type	Size	Value
batch_size	int	1	100
mnist	contrib.learn.python.learn.datasets.base.Datasets	3	Datasets object of...
n_classes	int	1	10
n_nodes_h1	int	1	500
n_nodes_h2	int	1	500
n_nodes_h3	int	1	500

The IPython console at the bottom shows the output of the script, indicating that the training process has completed 10 epochs with a decreasing loss and an accuracy of 0.9511.

```
See 'tf.nn.softmax_cross_entropy_with_logits_v2'.

Epoch 0 completed out of 10 loss: 1666037.4677734375
Epoch 1 completed out of 10 loss: 377809.3128890991
Epoch 2 completed out of 10 loss: 201302.4857263565
Epoch 3 completed out of 10 loss: 119427.91378033161
Epoch 4 completed out of 10 loss: 72651.25679710507
Epoch 5 completed out of 10 loss: 45327.621502393486
Epoch 6 completed out of 10 loss: 31955.17812934518
Epoch 7 completed out of 10 loss: 23664.35610633137
Epoch 8 completed out of 10 loss: 18248.740643078025
Epoch 9 completed out of 10 loss: 19962.00065876091
Accuracy: 0.9511

In [2]:
```